

Generalizing Treatment Effect Estimates From Sample to Population: A Case Study in the Difficulties of Finding Sufficient Data

Evaluation Review
2017, Vol. 41(4) 357-388
© The Author(s) 2016
Reprints and permission:
sagepub.com/journalsPermissions.nav
DOI: 10.1177/0193841X16660663
journals.sagepub.com/home/erx



Elizabeth A. Stuart^{1,2,3} and Anna Rhodes⁴

Abstract

Background: Given increasing concerns about the relevance of research to policy and practice, there is growing interest in assessing and enhancing the external validity of randomized trials: determining how useful a given randomized trial is for informing a policy question for a specific target population. **Objectives:** This article highlights recent advances in assessing and enhancing external validity, with a focus on the data needed to make ex post statistical adjustments to enhance the applicability of experimental

¹ Department of Mental Health, Johns Hopkins Bloomberg School of Public Health, Baltimore, MD, USA

² Department of Biostatistics, Johns Hopkins Bloomberg School of Public Health, Baltimore, MD, USA

³ Department of Health Policy and Management, Johns Hopkins Bloomberg School of Public Health, Baltimore, MD, USA

⁴ Department of Sociology, Johns Hopkins University, Baltimore, MD, USA

Corresponding Author:

Elizabeth A. Stuart, Department of Mental Health, Department of Biostatistics, Department of Health Policy and Management, Johns Hopkins Bloomberg School of Public Health, 624 North Broadway, Room 839, Baltimore, MD 21205, USA.

Email: estuart@jhu.edu

findings to populations potentially different from their study sample.

Research design: We use a case study to illustrate how to generalize treatment effect estimates from a randomized trial sample to a target population, in particular comparing the sample of children in a randomized trial of a supplemental program for Head Start centers (the Research-Based, Developmentally Informed study) to the national population of children eligible for Head Start, as represented in the Head Start Impact Study.

Results: For this case study, common data elements between the trial sample and population were limited, making reliable generalization from the trial sample to the population challenging. **Conclusions:** To answer important questions about external validity, more publicly available data are needed. In addition, future studies should make an effort to collect measures similar to those in other data sets. Measure comparability between population data sets and randomized trials that use samples of convenience will greatly enhance the range of research and policy relevant questions that can be answered.

Keywords

causal inference, generalizability, Head Start Impact Study, REDI evaluation, transportability

Randomized controlled trials remain the “gold standard” for research designs to estimate the effects of interventions. However, a shortcoming of nearly all randomized trials across a variety of fields (including medicine, public health, education, and social program evaluation) is that their measured effects are formally only generalizable to the subjects (e.g., students, patients, service providers) within the trial itself, leaving open the question of whether the intervention would be effective in a different target population. Important policy questions regarding program implementation may involve populations quite different from those in the trial, in which case practitioners and stakeholders must evaluate how well an existing trial can inform a particular policy decision. This issue is of particular relevance given increasing concerns about gaps between research and practice. Recent evidence across a number of fields has indicated that the subjects in a trial are often quite different from the individuals of policy or practice interest (e.g., Humphreys, Weingardt, & Harris, 2007; Rothwell, 2005; Stirman, Derubeis, Crits-Christoph, & Rothman, 2005; Westen, Stirman, & DeRubeis, 2006). The topic has received particular attention in mental health. For

example, Braslow et al. (2005) found that minorities were often underrepresented in studies of psychiatric treatment, and Susukida, Crum, Stuart, and Mojtabai (2016) found that, on average, individuals who participate in randomized trials of drug abuse treatment have higher levels of education and are more likely to be employed than are individuals seeking drug abuse treatment nationwide. In education, Bell, Olsen, Orr, and Stuart (2016) have shown that these differences can translate into bias when estimating population effects.

More formally, randomized trials offer internal validity: unbiased estimation of treatment effects in the sample of individuals in the trial but do not necessarily offer external validity: “whether the causal relationship holds over variation in persons, settings, treatment, and measurement variables” (Shadish, Cook, & Campbell, 2002, p. 20). Bareinboim and Pearl (2013) define a closely related term, “transportability,” as “a license to transfer causal effects learned in experimental studies to a new population, in which only observational studies can be conducted” (p. 107). Cook (2014) gives a broad overview of the challenges in generalizing from one sample to another, including discussion of extrapolation to new settings and contexts. What we are concerned with in this article is how we can estimate the average treatment effect in a target population of interest, given data from a randomized trial. In other words, how can policy makers and practitioners who want to implement best practices in their programs, schools, and clinics use existing randomized control trial evidence to evaluate the effects of an intervention in their target population? How feasible for practitioners are methods for evaluating the effects of existing randomized control trials in different populations of interest?

In this process, defining the target population is a crucial step, and the choice of target population will depend on the specific policy or practice question of interest; for some questions, the appropriate target population may be a national population, while for others, it may be a more narrow subpopulation, such as participants in a federal program in a particular state. (And in fact, the same randomized trial may be used to generalize effects to multiple target populations.) Throughout this article, we assume that the target population is well defined and appropriate for the question of interest.

Fundamentally, problems of external validity arise when there are factors that moderate treatment effects and that differ between the trial and the population (Cole & Stuart, 2010; Olsen, Bell, Orr, & Stuart, 2013). If the composition of subjects in the trial differs from that in the population of

interest on effect moderators, then the average effect in the trial may not reflect the average that would be observed in the population. Given increasing concerns about external validity (see, e.g., Orr, 2015), statisticians and methodologists have begun to develop statistical methods to assess and enhance external validity. An ideal research process for improving external validity is to begin by formally defining a target population and then carrying out a randomized trial in a sample selected to be representative of that population. However, this sample selection process is generally difficult to carry out and quite rare; Olsen, Bell, Orr, and Stuart (2013) document that only 7 of the 273 studies in the *Digest of Social Experiments* (Greenburg & Schroder, 2004) were conducted in random samples from the population of interest. (And in fact, in many cases, the target population itself is also not well defined.) Additional approaches for assessing and enhancing external validity utilize multiple studies on the same intervention (e.g., individual patient meta-analysis that combines multiple randomized trials, or cross-design synthesis, which combines experimental and nonexperimental evidence; see Stuart, Bradshaw, & Leaf, 2015, for an overview of these design and analysis approaches).

The turn toward an emphasis on external validity is driven by an interest in using the results of internally valid randomized controlled trials to inform practice in other settings and with different target populations. Other papers (as cited above) introduce and illustrate statistical methods for assessing and enhancing generalizability; the current article examines the practical issues that arise in the process of employing one of the proposed methods for estimating target population treatment effects using a randomized trial sample and population data. In this article, we focus on settings where a single randomized trial has been conducted, and where there is interest in generalizing the results of that trial to a specific target population, but where we do not have information in the population about possible treatment receipt (only covariates and possibly outcomes are available). Some recent work (e.g., Cole & Stuart, 2010; O'Muircheartaigh & Hedges, 2014; Stuart, Cole, Bradshaw, & Leaf, 2011) has proposed reweighting approaches that weight the trial sample to look like the target population on a set of key characteristics. However, there has been relatively little use of those methods and little investigation of how well they perform, or how feasible they are in practice. This article focuses on the assumptions and data needed to utilize a reweighting method to estimate treatment effects in the target populations and highlights recommendations for the design of future studies and data sets to help facilitate such estimation.

Method

Formal Setting for External Validity

We first clearly define the quantity of interest using the concept of potential outcomes (Rubin, 1974). Consider a setting where we are interested in estimating the effect of a supplemental enrichment program for children in Head Start. After identifying the treatment and comparison conditions of interest, the next step is clearly defining the target population of interest. For our case study, detailed further below, the target population is the entire population of students affected by Head Start policies: all children in Head Start programs across the United States. Each individual in the population of interest has two potential outcomes: their outcome (e.g., a measure of emergent literacy skill scores) if they receive the intervention of interest (e.g., the supplemental enrichment program), denoted Y_{1i} , and their outcome (emergent literacy skill scores) if they receive the comparison condition of interest (e.g., standard instruction), denoted Y_{0i} . The treatment effect for individual i is defined as the difference between these two potential outcomes: $\Delta_i = Y_{1i} - Y_{0i}$. In this article, our interest is in estimating the population average treatment effect (PATE), denoted Δ : the average Δ_i across all N individuals in the target population:

$$\Delta = \frac{1}{N} \sum_{i=1}^N \Delta_i.$$

In this case study, we utilize data from an existing randomized trial of a supplemental enrichment program, the Research-Based, Developmentally Informed (REDI) program, in which children in Head Start centers were randomized to receive the supplemental program or not. Let P_i denote membership in the population of interest ($P_i = 1$ for children in the population and $P_i = 0$ for those in the trial sample). We assume that the same specific individuals are not in both the sample and the population, although a slight modification of the weights defined below allows for the individuals in the sample to be a subset of the individuals in the population data set, as in Stuart, Cole, Bradshaw, and Leaf (2011). Let T_i denote treatment assignment in the trial ($T_i = 1$ for intervention and $T_i = 0$ for control). If treatment assignment is randomized, the difference in potential outcomes between the observed treatment and control groups in the trial sample (those for whom $P_i = 0$), $\hat{\Delta} = \bar{Y}_{P=0, T=1} - \bar{Y}_{P=0, T=0}$ (where $\bar{Y}_{P=p, T=t}$ denotes the sample mean of Y for individuals in population p and treatment condition t), provides an unbiased estimate of the treatment effect in the trial sample but

may not provide a good estimate of the PATE. We define the external validity bias as the difference between the impact estimated in the evaluation sample and the true population impact, $\hat{\Delta} - \Delta$.

Imai, King, and Stuart (2008) provide a framework for thinking about this bias, decomposing the overall bias when estimating a population treatment effect using an experimental or nonexperimental study sample into four pieces: internal validity bias due to observed characteristics, internal validity bias due to unobserved characteristics, external validity bias due to observed characteristics, and external validity bias due to unobserved characteristics. When interest is in estimating the PATE, standard beliefs about a randomized trial being the gold standard design may not hold, if, for example, the external validity bias of a small, non representative trial is larger than the internal validity bias of a large-scale non-experimental study conducted in a representative sample. This article focuses on trying to reduce the bias due to observed differences between the sample and population. Cole and Stuart (2010) and Olsen et al. (2013) provide analytic expressions for external validity bias as a function of the extent of treatment effect heterogeneity and the differences between a trial sample and the target population of interest.

Measures for Assessing External Validity

Stuart et al. (2011) and Tipton (2014) propose indices that summarize the similarity between a randomized sample and a population. Stuart et al.'s (2011) index is based on the mean difference between the predicted probabilities of participating in the trial, compared between the units in the trial and those in the target population (similar to the idea of assessing propensity score overlap in a nonexperimental study). Tipton's (2014) index is based on the Bhattacharyya coefficient (1943), which measures the similarity of the densities of predicted logit probabilities between the trial and population. Tipton (2014) provides some rules of thumb for interpreting the index as a measure of the similarity between the trial sample and the population and for when the trial is sufficiently similar to the population to enable generalization. In the Estimation of Weights and Covariate Diagnostics section, we illustrate the calculation of these metrics in the case study for this article.

Reweighting Approach to Estimate Population Treatment Effects

In this article, we consider a reweighting approach for equating a randomized trial sample and a target population, with the goal of estimating the

target PATE. It is useful for settings where there is interest in generalizing from a trial to a target population, and the only available data are data from a single trial and covariate data on the population. (If there are multiple trials conducted, other approaches, such as cross-design synthesis, may be more useful.) While in this article we focus on the reweighting approach, other model-based approaches such as Bayesian Additive Regression Trees (BART; Kern, Stuart, Hill, & Green, 2016) could be used instead and would generally have the same data needs as the reweighting approach.

The reweighting approach is related to inverse probability of treatment weighting for estimating causal effects in nonexperimental studies and nonresponse weighting adjustments for handling survey nonresponse. The main idea is to weight the randomized trial sample to look like the population of interest on a set of key covariates. The specific procedure is as follows:

1. Create a combined (stacked) data set with the randomized trial sample and the population data set, with a set of covariates X observed in both groups.
2. Create an indicator variable for being in the target population (P).
3. Estimate a model of membership in the population (P) as a function of the covariates X , for example, using logistic regression.
4. Create weights for individuals in the trial sample, defined as $w_i = \frac{p_i}{(1-p_i)}$, where $p_i = P(P_i = 1|X_i)$.
5. Estimate treatment effects using the individuals in the randomized trial by running a weighted regression of outcome as a function of treatment status and the covariates, with the weights calculated in Step 4.

Details of this process using the case study are provided below.

There are three structural assumptions underlying the reweighting approach investigated (see Stuart et al., 2011, for more details):

Assumption A1: Given the observed covariates X , every subject in the population has a nonzero probability of participating in the randomized trial. If this were not true, there would be some set of individuals in the population who would not be represented at all in the trial; generalizing to them would require extrapolating outside the range of the data in the trial.

Assumption A2: Unconfounded sample selection: There are no unobserved variables related to selection into the trial and

treatment effects, given the observed covariates X . The implication of this assumption is that we need to observe and adjust for all factors that drive selection into the trial and moderate treatment effects. However, we do not need to adjust for factors that relate to only one of these mechanisms (e.g., a factor that is related to inclusion in the trial but that does not moderate treatment effects).

Assumption A3: Treatment assignment is random in the trial.

These are important assumptions to consider in any empirical example and are discussed further in the context of the case study below.

Data Needs for Using These Approaches

In order to utilize the reweighting approach, one must have data on both the randomized trial of an intervention of interest and a specific target population. The necessary data elements from the randomized trial sample include a treatment indicator (treatment vs. control), covariates, and the outcome(s) of interest. In the target population data set, one needs covariate information. Although some approaches, such as BART, can utilize outcome data from the population, the reweighting approach focused on in this article does not use such data; it only requires covariate data in the population. (And in fact if outcome data on the population is available, methods using BART may be preferable; Kern et al., 2016.) If outcome data under the control condition are available in the population (e.g., if the control condition in the trial was “usual care” and no one in the population received the treatment of interest), that data can also be useful as a check on the similarity between the weighted trial sample and the population (Stuart et al., 2011).

For the reweighting approach, covariates within the sample and population data sets must possess sufficient overlap to allow for the generation of a group of “common” covariates that can be used to estimate membership in the target population and that will plausibly satisfy Assumption A2; in particular, the reweighting approach can only adjust the randomized trial to look like the population with respect to observed characteristics. If the trial and population differ on an unobserved variable (or variables) that moderates treatment effects (a violation of Assumption A2), the estimated PATE will be biased (except in pathological cases where biases caused by different variables may cancel each other out). Because of this, it is generally appropriate to include as many covariates in the set of covariates X as possible; this is similar to the advice in the propensity score literature that it is generally better to err on the side of including covariates rather than

excluding them when using propensity scores to estimate treatment effects in nonexperimental studies (Myers et al., 2011; Stuart, 2010). It is especially crucial to include likely effect moderators, either based on data analysis of the randomized trial data (e.g., subgroup analyses or other methods for assessing treatment effect heterogeneity) or through a conceptual model of impacts (e.g., which factors program developers think may moderate treatment effects; that conceptual model would also presumably influence the potential moderators measured and examined in the trial). For example, a particular intervention may be hypothesized (or observed in the trial) to be more effective for individuals with particularly low baseline levels of achievement, in which case it would be important to observe those baseline achievement levels in the trial and population.

Background and Selection of the Case Study

The overall goal of this article, and thus the selection of a motivating case study, is to illustrate the feasibility of using these newly developed statistical methods for estimating population treatment effects using existing data. A clean and “well-behaved” case study would be one that has a randomized trial conducted with a clear and well-defined target population, and a comprehensive set of covariates available for both the trial and the target population; this is the type of case study often used in papers describing new statistical methods. In contrast, the case study detailed below was not selected because it is particularly well behaved or ideal for the statistical method, but rather to focus on the practical implications and challenges in using the generalizability method in the real world. The growing focus on external validity (Orr, 2015) demands practical examples to illustrate the weaknesses and strengths of our current methods of data collection and statistical methodologies for the purposes of evaluating external validity and estimating treatment effects in populations outside trial-specific samples.

To illustrate the use of these methods for estimating population treatment effects, and their complications in practice, we chose to select a case study in the area of early childhood education. This choice was driven by a growing interest in methods to assess and enhance generalizability in that field (e.g., as exemplified by a 2014 meeting on external validity organized by the Office of Planning, Research, and Evaluation within the Administration for Children & Families of the U.S. Department of Health and Human Services), and the lack of examples of their use in that area. Within the broad field of early childhood education, we chose to focus on

interventions that evaluated academic outcomes in order to identify a randomized trial and a target population to illustrate the statistical methods.

Even though significant resources have been invested in making information from randomized control trials in education available to practitioners through resources such as the What Works Clearinghouse (WWC) and Institute for Education Sciences websites, a search of these and other online sources (detailed in Appendix) produced few available options for access to randomized trial sample and target population data. The Appendix details the process we used, and some of the challenges we encountered in identifying a case study, particularly a lack of public use data from randomized trials and a lack of comparable measures across data sources.

The randomized trial we ultimately identified for our case study was of the Head Start REDI intervention, which involved an enrichment program that included lessons, extension activities, teaching strategies, and teacher support. The enrichment was randomly assigned to Head Start classrooms: 44 classrooms were assigned to the intervention or to a control group that consisted of maintaining “usual practice.” The intervention took place over the course of 1 year in Head Start for 356 four-year-olds. The Head Start classrooms that participated in the REDI study are all located in three counties in Pennsylvania. The outcomes of interest for the REDI study were language development, emergent literacy, and social-emotional competencies. The battery of assessments included child assessments, teacher ratings, parent ratings, and direct classroom observations (Bierman et al., 2008). The study found significant effects of the intervention on 7 of the 11 language, emergent literacy, emotional understanding, and social problem-solving skill measures. For our case study, we chose to focus on a single academic outcome, the Elision test, a measure of phonological awareness and an indicator of emergent literacy skills, which showed significant positive effects in the REDI trial.

Selection of the Target Population

In a practical application of the reweighting method, a practitioner or policy maker will have a target population in mind, for which they are interested in evaluating the potential impact of a given intervention. For our case study, we are interested in determining how well the REDI intervention would work for Head Start students across the country. To evaluate the impact of the intervention in a target population, one must obtain data on that population of interest. For the REDI case study, two main data sets provide the potential for population level, early childhood data: the Early Childhood

Longitudinal Birth Cohort (ECLS-B) Study and the Head Start Impact Study (HSIS). Since the REDI trial was conducted in Head Start centers, we chose the HSIS as our target population data set to illustrate the reweighting method. The positive impact from the REDI trial raises the question of whether similar effects would be observed if the intervention was more broadly implemented in Head Start classrooms nationally. The program evaluators argue that their findings of a positive effect of the intervention on students' literacy skills suggest that, "it is possible to integrate empirically validated strategies for promoting these critical emergent literacy skills in ways that are consistent with Head Start practices" (Bierman et al., 2008, p. 1813). We thus used the nationally representative HSIS data set available through the Inter-university Consortium for Political and Social Research (ICPSR; <http://doi.org/10.3886/ICPSR29462.v5>) to evaluate the impact of integrating the REDI intervention in the national Head Start population.

The HSIS is a nationally representative sample of Head Start programs and children (U.S. Department of Health and Human Services, 2010). In fact, the HSIS is itself a randomized controlled trial designed to evaluate the effects of Head Start; it was carried out among centers with waiting lists (to facilitate random assignment) and thus results are formally representative of the population of children in centers with waiting lists. Children who applied for Head Start were randomly assigned, in separate 3-year-old and 4-year-old cohorts, to the treatment group, which was allowed to enroll in Head Start, or a control group, which could not enroll in Head Start that year. Using the HSIS data allows us to evaluate the impacts of REDI in the target population of Head Start-eligible students; however, other policy makers or practitioners might be interested in the effects of the REDI program in a different target population. For example, there may be interest in the effects of the REDI intervention on all 4-year-olds, including a more socioeconomically diverse population of children than those eligible for Head Start, in which case a nationally representative data set of children in the United States, such as the ECLS-B, would be a more appropriate population data set (<http://nces.ed.gov/ecls/birth.asp>).

Results

This section details the use of the REDI and HSIS data to evaluate the generalizability of the REDI program to Head Start-eligible children nationwide, as represented by the HSIS. Given the HSIS sampling, we work from the specific motivating question, "What would the impacts of the

REDI intervention have been for children in all Head Start centers in the nation (with waiting lists) in approximately 2002?”

For our analyses, we use both the REDI randomized controlled trial sample of 4-year-olds enrolled in Head Start centers in Pennsylvania ($N = 352$) and the nationally representative sample of Head Start–eligible 4-year-olds in the HSIS ($N = 1,983$).¹ For the analyses, we utilize the HSIS child base weight, which weights the HSIS sample to be nationally representative. For simplicity in this illustrative case study, we will treat the children as the units of analysis (rather than Head Start classrooms or centers); calculation of standard errors accounts for clustering at the center level. To address missing data on covariates and outcomes, we ran a single imputation in each data set separately (REDI and HSIS) and then appended the data sets to generate the common covariates (detailed below). The largest missing values in the REDI data set were for mother’s education and mother’s marital status, both with 11% missing values. The Elision outcome variable had 6% missing values, while other common covariates had less than 3% missing values in the REDI data set.

Measures

Outcomes. The outcome of interest examined was the Elision assessment of phonological processing from the Test of Preschool Early Literacy.² The Elision test measures the ability to remove phonological segments from spoken words to form other words.³ Preschoolers who possess this and other early phonological awareness skills are more proficient in reading skills during first and second grade, even after controlling for vocabulary skills and student IQ (Bryant, MacLean, Bradley, & Crossland, 1990; Catts, Fey, Zhang, & Tomblin, 1999). The Elision test showed a significant intervention effect in the REDI program evaluation.

Common covariate measures. In addition to the outcome measure, we were able to create a total of seven covariates measured comparably across the two data sets: male, race/ethnicity (Black, Hispanic, and White/Other), household size, mother’s marital status, mother’s education level, a baseline measure of the Applied Problems mathematics test from the Woodcock Johnson III, and a baseline measure of the Elision test score.⁴ Some were easily defined, such as child gender, while others we were able to make comparable by combining or altering existing variables in the appropriate data sets; full details on how they were made comparable across the two data sets is provided in Appendix B.

As detailed in Appendix B, three covariates required significant editing to generate variable structures that could be applied to both REDI and HSIS, including race, mother's marital status, and mother's education. The difficulty in generating common variables for such simple demographic measures demonstrates the importance of survey data collection processes for future applications of the proposed reweighting method and other post hoc analyses to assess generalizability. When survey items are designed differently, or when surveys collect different types of demographic information, generating common variables can be difficult.

Additionally, some similar constructs in both the REDI and HSIS data sets were measured using different tests or assessments, such that they could not be used to generate common covariates. For example, the HSIS used the Peabody Picture Vocabulary Test, while REDI used the Expressive One-Word Picture Vocabulary Test. Although both assess an ability to use words to describe pictures, the use of different tests made it difficult to make these variables comparable, and ultimately they were excluded from the model.

Estimation of Weights and Covariate Diagnostics

As a reminder, the generalizability method illustrated in this article aims to reweight the trial (REDI) sample to look like the target population (HSIS) in order to equate the two samples with respect to a set of observed covariates. To do this, the reweighting method begins with a logistic regression model to estimate the probability of membership in the HSIS target population. The dependent variable is membership in the HSIS population ($P_i = 1$ for individuals in the HSIS target population; $P_i = 0$ for individuals in the REDI trial sample); the data is weighted by the HSIS base weight to reflect the population. Predictors in the model were variables for gender, race, family size, mother's marital status, mother's highest level of education, a baseline measure of the Applied Problems mathematics test, and a baseline measure of the Elision test score. The predicted probabilities from this logistic regression model were then used to calculate the weights for the REDI sample, using the equation in the Method section. This weight is used to weight the REDI subjects to resemble the nationally representative HSIS target population; see Kern, Stuart, Hill, and Green (2016) and Hirano, Imbens, and Ridder (2003) for more details on this weighting, which is analogous to "weighting by the odds" when estimating the average treatment effect on the treated in nonexperimental studies.

Before describing the weights themselves, we present the Tipton (2014) and the Stuart et al. (2011) measures of similarity between the REDI




Figure 1. Distribution of logit propensity scores in the REDI sample and HSIS population. REDI = Research-Based, Developmentally Informed; HSIS = Head Start Impact Study.

randomized trial sample and the population of Head Start–eligible students. In this case study, the Tipton metric was .93, a value that implies “very high” generalizability according to Tipton (2014). This interpretation indicates that the REDI sample is like a random sample drawn from the broader Head Start population, at least with respect to the observed covariates. Stuart et al. (2011) metric produces a less optimistic evaluation of the overlap between the sample and the target population, producing an absolute standardized difference in means between the propensity scores for the sample and population of .73. This indicates that differences between the sample and population may be large enough to result in unreliable estimates due to extrapolation. The difference between the Tipton and Stuart metrics is potentially consistent with results in Tipton, Hellberg et al. (In Press), showing that simply by chance the standardized mean difference (SMD) can be large in small and moderate random samples. Another explanation could potentially be that given the small size of the study sample relative to the size of the population, the average propensity scores are all quite large and clustered around 1 (Figure 1). However, the different results for the two metrics indicate the need for more work to

understand the differences between these metrics and the implications for our conceptualization of external validity.

The resulting weight included some large values, as can happen with inverse weighting (Schafer & Kang, 2008). To limit the influence of extreme weights, the weights were trimmed at the 95th percentile, with all values above the 95th percentile set equal to the value at the 95th percentile (as in Lee, Lessler, & Stuart, 2010). In particular, 18 cases had their weights truncated to approximately 3,000 (the 95th percentile); the largest weight before trimming was approximately 8,000, but most trimmed weights were closer to 4,500 before trimming. (Note that the mean of the weights is also large, due to the large size of the HSIS population relative to the REDI sample.) It is important to note that a set of individuals with large outlying weights could indicate large differences between the trial sample and the target population, and possibly segments of the population that are not well represented in the trial. This can also be identified through density plots of the propensity scores themselves, as shown in Figure 1. If large areas of nonoverlap are found, researchers may need to refine the definition of the target population to reflect a subset for which more reliable generalizations can be obtained; see Tipton, Feller et al. (in press) for a strategy for doing so.

In addition to examining the weights, another diagnostic, the SMD, can be used to compare unweighted and weighted covariate means between the trial sample and population. The SMD is calculated as the difference in means between groups (e.g., unweighted REDI sample and HSIS population) divided by the standard deviation of the pooled values; if the weighting is successful, the samples should look more similar after the weights are applied. Table 1 compares the REDI and HSIS samples with respect to the seven commonly measured covariates, before and after the propensity score weighting. In Table 1, "REDI unwt." refers to the original REDI sample, without the propensity score weights. "REDI wtd." refers to the REDI sample, but with the propensity score weights applied. Similarly, the SMD columns indicate "unwtd." and "wtd." to reflect without and with the propensity score weights, respectively. All HSIS calculations (e.g., the HSIS mean column) use the HSIS base weights to ensure national representativeness, as denoted by the "base wt." notation.

The SMD calculations shown in Table 1 generally indicate that there is improved covariate balance after the weighting; most variables show a smaller SMD following the weighting indicating that the REDI sample looks more similar to the HSIS population. There are a few exceptions, including gender, mother having some postsecondary education, and Black that had mean values after weighting that were slightly further from the

Table 1. Comparison of REDI and HSIS Samples.^a

Variables	REDI Unwtd. Mean	HSIS Base Wt. Mean	REDI Wtd. Mean	Std. Mean Difference Unwtd.	Std. Mean Difference Wtd.
Males	0.46	0.51	0.44	.1	.13
White/Other	0.66	0.32	0.38	-.71	-.12
Black	0.16	0.18	0.21	.06	-.06
Hispanic	0.19	0.5	0.42	.63	.16
Family size	4.59	4.93	4.69	.19	.14
Mother married	0.36	0.5	0.42	.29	.16
Mother high school education	0.42	0.32	0.31	-.22	.02
Mother postsecondary education	0.28	0.26	0.29	-.07	-.09
Applied standard score pretest	94.88	87.04	89.44	-.48	-.15
Elision pretest	8	6.79	7.23	-.34	-.12

Note. HSIS = Head Start Impact Study; REDI = Research-Based, Developmentally Informed.
^aThe column labeled "base wt." use the Head Start Impact Study weight that ensures the mean is nationally representative. "unwtd." refers to no weighting. The "std. mean difference wtd." column refers to weighting by the generalizability population weights defined in the second section.

HSIS mean as compared with their unweighted means. Although propensity score theory says that in large samples all covariates should see more similarity following the weighting (Rubin & Thomas, 1996), in real data (especially small samples), it is not uncommon for some variables to become more different following the propensity score adjustment. This is particularly common for variables that did not show large differences before the weighting, as we see in Table 1; gender, Black, and mother having some postsecondary education had the smallest SMDs before the weighting (and still relatively small SMDs following the weighting). This pattern occurs because the variables showing large differences between the sample and population are the ones that drive the propensity score model (for further discussion, see Stuart, Lee, & Leacy, 2013).

Estimating the SATE and PATE

In the full evaluation of the REDI intervention, the Elision test was positively and significantly affected by the intervention, with an effect size of

.35, $p = .001$. In the program evaluation, the intervention effect was modeled using a hierarchical linear model, with child race, child gender, and a preintervention score as Level 1 covariates and site (central or Southeastern Pennsylvania), cohort, and intervention status as Level 2 covariates (Bierman et al., 2008).

For our case study, we can only replicate the REDI evaluation model to a certain extent. We use the common covariates we were able to generate to run a simple regression, using the covariates that were also in Level 1 of their hierarchical model: child gender, child race, and the pretest score for the Elision test. For the analysis, we run two models. First, an unweighted model, to estimate the SATE: the effect of the intervention on the Elision test in the randomized controlled trial sample. Second, a weighted model that uses the weights defined above to estimate the PATE: the effect of the REDI intervention in the target population, the nationally representative population of Head Start–eligible 4-year-olds. The standard error calculations use survey methods (Taylor series linearization) to account for the weighting and clustering (see, e.g., McCaffrey, Ridgeway, & Morral, 2004).

As a reminder, interpreting the results of the weighted model as an estimate of the PATE requires the three assumptions detailed above:

Assumption A1: Given the observed covariates X , every subject in the population has a nonzero probability of participating in the randomized trial. The density plot and SMD diagnostics make this assumption seem reasonable in these data, in that generally good covariate balance is obtained, and there is overlap across the range of propensity scores.

Assumption A2: Unconfounded sample selection: There are no unobserved variables related to selection into the trial and treatment effects, given the observed covariates X . This assumption is likely more questionable in this example, especially given the difficulty in finding common measures. For example, it is plausible that family income would moderate the effects of the program, but this measure is not available for adjustment. Similarly, effects could vary across levels of cognition or other measures of reading or other skills not captured by the Applied Problems standard score pretest or the Elision pretest.

Assumption A3: Treatment assignment is random in the trial. This assumption is satisfied, given the conduct of the HSIS.

Table 2. Estimates of the Sample Average Treatment Effect (SATE) and Population Average Treatment Effect (PATE).

Variables	SATE (SE)	<i>p</i> Value	PATE (SE)	<i>p</i> Value
Treatment	0.66 (0.45)	.16	1.11 (0.54)	.05
Constant	9.34 (0.70)	.00	9.56 (0.89)	.00
N	352		352	

Table 2 shows the estimates of the SATE and PATE. The SATE estimate is not statistically significant at the $\alpha = .05$ level. However, the PATE estimate is larger and marginally significant with a p value = .05. This indicates that the REDI intervention, if implemented for the entire target population of Head Start–eligible 4-year-olds, would have a positive and marginally significant effect on students’ Elision test, an indicator of phonological awareness, if Assumptions A1 to A3 are met in this analysis. It should be noted that the more sophisticated estimation techniques used in the original evaluation of the REDI sample do find significant effects for the treatment on students’ outcomes in the randomized controlled trial sample. The more simplistic case study analysis provides some evidence that these positive findings would be replicated if the intervention was scaled up and expanded to all Head Start centers nationally. Although in this case study similar conclusions are obtained with respect to the SATE and PATE, this is not guaranteed and other examples may see large differences between the SATE and PATE.

One note is that the standard errors of the PATE estimate are larger than those of the SATE estimate, which is common when trying to generalize results from sample to population, in part because of the extrapolation (and thus uncertainty) inherent in doing that generalization. Larger weights, indicating more extrapolation, will make the increase in standard errors even greater.

It is important to note that although the common covariates we generated allow for an estimation of the treatment effect on the Elision posttest score in some capacity, the number of common covariates we were able to produce remains small, limiting our ability to develop the best possible model for generating a weight for the REDI sample to use in these analyses, as discussed further below. However, our model serves as a basic illustration of the reweighting methodology for assessing the generalizability of the effect estimate for the Elision test in the REDI randomized controlled trial, and we find the potential for a positive effect of the REDI intervention in a broader target population.

Evaluating Effect Heterogeneity

As discussed earlier, the reason why the SATE and PATE may differ is if there are variables that moderate treatment effects and that differ between the sample and population. Thus, methods to detect treatment effect heterogeneity are inherently related to methods that assess generalizability. To develop a greater understanding of why the PATE estimate may be larger than the SATE estimate in the motivating example, we (post hoc) examine whether the variables with large differences between the REDI sample and the HSIS target population moderate treatment effects. In particular, we investigated whether there is evidence of any effect heterogeneity across racial and ethnic subgroups (White, Black, and Hispanic) or across levels of the Applied Problems standard score. The Applied Problems test is a math test that assesses students' quantitative reasoning and math knowledge. Completing this test requires students to construct mental math models through the use of their language comprehension, calculation, and math skills (Wendling, Schrank, & Schmitt, 2007). This was one variable that showed the largest SMDs between the trial sample and the population (Table 1).

To investigate potential effect heterogeneity across these variables, we estimated impacts in the (unweighted) REDI sample using three separate models to include interaction terms for Black, Hispanic, and the Applied Problems standard score. The models also included the same covariates from the SATE and PATE estimates: child gender, child race, and pretest score for the Elision test. For the Applied Problems standard score, we included the Applied Problems pretest in the model and the interaction term of Applied Problems with treatment. We test for effect heterogeneity because we might hypothesize, for example, that due to lower scores on the Applied Problems pretest in the HSIS population (mean = 87.0), on average, than in the REDI sample (mean = 94.9); weighting the REDI sample to look like the HSIS population could make the impact go up, since individuals with higher impacts (lower applied problems test scores) are more represented in the HSIS population than in the REDI sample.

These analyses, however, demonstrated no evidence of effect heterogeneity across the Applied Problems score or the racial and ethnic categories (Table 3); none of the interaction terms reach statistical significance.

Conclusions drawn from subgroup analyses should be interpreted cautiously, given that these analyses are beset by the challenges in all randomized trials (Supplee, Kelly, MacKinnon, & Yoches Barofsky, 2013), in particular worries about multiple comparisons and concerns about limited

Table 3. Interaction Terms From Three Models Predicting Elision Scores and Examining Effect Moderation.

Interaction Term	Coefficient (SE)	p Value
Treatment × Black	0.14 (1.06)	.894
Treatment × Hispanic	0.87 (1.00)	.395
Treatment × Applied Problems Standard Score	−0.06 (0.03)	.099

power to detect effect moderators. However, evaluating the possibility of effect heterogeneity is an important step in the analysis process because these analyses may help explain differences between PATE and SATE estimates.

Discussion

This article has provided a case study of using an existing randomized controlled trial, combined with data from a target population of interest, to estimate the PATE of a treatment condition of interest. Although new statistical methods are promising for estimating population treatment effects using existing data, this article shows that the feasibility of these methods in more general practice depends on (1) generating relevant population data sets and (2) ensuring measure comparability between trials and population data sets.

Data availability is crucial for answering questions about generalizability. In particular, data are needed on randomized trials and on target populations for all the important moderators of intervention impact. The data gathering process for this case study involved multiple steps, including evaluating studies for high-quality randomized controlled trials and accessing restricted data through license applications or contacting investigators. The small number of publicly available high-quality randomized controlled trials with individual-level data available is a significant limiting factor to generalizability assessments for policy makers and practitioners. There are many high-quality trials that have evaluated interventions in early childhood education, and the results of these studies are collected and made available to researchers and practitioners through tools such as the WWC and the Institute for Education Sciences websites. As a result, practitioners generally only have access to results valid for the participants in the randomized controlled trial, without the tools necessary to evaluate the external validity of these results and the tools to evaluate the effects of these trials in

their particular population of interest. Some data are available through restricted access, including the HSIS used in our case study, and in these cases, the data licensing step serves as an important component of the data procedures to protecting respondent confidentiality. However, many studies (even those funded through federal grants which frequently include policies for making data publicly available) do not have procedures in place that would allow practitioners or researchers to apply for access to these data sets. Greater access to individual level data from these studies would enhance our abilities to use them to answer questions that may go somewhat beyond the specific aims of the original studies. We note that in K–12 education, there is often school-level data available (e.g., through the Common Core of Data), which can be used to establish populations of interest, but such data are still often limited in terms of the measures available, which could reduce confidence in Assumption A2 (no unmeasured effect moderators).

In addition, even when data sets exist and are accessible, many are extremely limited in terms of the overlap of key measures. The key underlying assumption of many of the existing methods for generalizing treatment effect estimates is that the variables that differ between the sample and population and that moderate treatment effects are observed. We had difficulty finding randomized trials and population data that had enough common covariates to make that assumption plausible. In the REDI and HSIS case study, we were able to utilize only seven commonly measured covariates, even though over 75 variables were available in the REDI data set and several hundred available in the HSIS data set. This allows us to demonstrate the reweighting method, but the set of covariates is insufficient to avoid concerns about the potential for bias in the PATE estimate. As a result of these concerns, the REDI case study presented here faces the problem of inadequate data to ensure reliable generalization from the sample to the population. The potential for evaluating the effects of local randomized controlled trials in population data sets in the future depends heavily on establishing common measures across population and trial data sets. Thus, in addition to making data more easily available, researchers should endeavor to collect standard (and standardized) measures in their studies, to help facilitate combining such data with other sources.

Of course, the measurement goals of a researcher carrying out a randomized trial are often quite different from the measurement goals of the designers of large population data sets; individuals carrying out an evaluation are often interested in whether impacts are seen on a measure specific to the core components and goals of the intervention. However, we encourage researchers and survey designers to consider the inclusion of a battery

of common measures to facilitate the combining of data sets to answer more complex and nuanced research questions. There is increasing interest in such integrative data analysis and data harmonization, and although there is some progress in developing methods that allow the combining of somewhat different measures (e.g., Bauer & Hussong, 2009), a set of common measures can go a long way in making the data more usable. As one model, some fields are moving toward a “common data model,” including initiatives such as the National Institutes of Health Toolbox (Weintraub et al., 2013), PROMIS (Cella et al., 2007), and PhenX (Hamilton et al., 2011). Orr (2015) argues that evaluators should adopt a two-stage model, with Stage 1 used to conduct a streamlined experimental evaluation and with a larger more complete evaluation at Stage 2 for those interventions that show particular promise in Stage 1. For studies that show promise in Stage 1, the inclusion of common measures in Stage 2 may help generate sufficient data for additional analyses of the generalizability of the effect estimates, expanding the potential use of the outcomes of a single trial for practitioners in different settings or serving different populations.

This key underlying assumption that we can adjust for all of the effect moderators (Assumption A2) also points to the need for more empirical and theoretical work understanding treatment effect heterogeneity. Our confidence in whether or not we can generalize results from a trial to a target population depends on whether we are confident that we have measured the relevant effect moderators. Although there have been recent advances in detecting effect heterogeneity (e.g., Kent, Rothwell, Ioannidis, Altman, & Hayward, 2010; Schochet, Puma, & Deke, 2014; Weiss, Bloom, & Brock, 2013), many studies are underpowered to detect effect heterogeneity and more work is needed in this area.

We also highlight that the reweighting approach for equating a randomized trial sample and a target population that we illustrate was developed only recently, and more work is needed to determine when this and other similar approaches work well, and how sensitive the results are to the underlying assumptions. Results in Kern et al. (2016) indicate that when the assumption of no unmeasured effect moderators is satisfied, flexible modeling approaches such as BART or reweighting approaches can work well, but when that assumption is not satisfied neither of the methods performs well. Further work should also consider how to extend these methods to multilevel settings to better account for the clustering of children within sites or schools. Additional work is also needed to evaluate the best methods for analyzing external validity and estimating generalizability under different conditions related to the underlying assumptions. To make these models

useful to policy makers and practitioners, clear guidelines about the best practice models under different conditions will be necessary.

Finally, a premise of the methods discussed here is that there is a well-defined target population. Any discussion of “generalizability” needs to be couched within the question “generalizability to whom?” In addition, a particular randomized trial may be used to generalize effects to multiple target populations (e.g., to individual states, for state-level decision-making, or to the nation as a whole), and a study may be generalizable to one population but not to another (see, e.g., Tipton, 2014).

In conclusion, as research studies in fields such as early childhood education become more and more rigorous in terms of their internal validity, there is growing interest in also assessing their potential external validity or generalizability. Researchers and policy makers would like to be able to answer more questions from existing data, for example, whether a policy maker can use the results from a given randomized trial to inform their decision-making for their population of interest (e.g., the Head Start director of a particular state). Statistical methods are beginning to be developed to estimate population treatment effects, but the data to use those methods appropriately are still lacking. That work is important and should continue. However, this article highlights, in the context of a very real-world example, that even with improving statistical methods, the research field is still far away from being able to confidently generalize results from randomized trials to target populations. And existing data limitations make it difficult to utilize these methods for practical evaluations of the effects of randomized controlled trials in populations of interest. Given the growing interest in generalizability, researchers conducting trials, and entities generating population data, need to incorporate considerations of external validity as part of their decision-making process about study design and data collection. Increased coordination within research fields, to move toward the collection of common measures, will greatly improve the potential for assessing external validity and generalizability, enhancing our ability to make informed decisions about the value of interventions in different settings and with different populations.

Appendix A

Identification of Case Study

To choose a recent randomized controlled trial for the case study, we established our area of interest as studies of early childhood

education with an academic outcome measure for literacy or math skills. To locate population data and a randomized trial, we utilized three main websites: the WWC (<http://ies.ed.gov/ncee/wwc/>), ICPSR (<https://www.icpsr.umich.edu/icpsrweb/landing.jsp>), and Childcare & Early Education Research Connections (<http://www.researchconnections.org/childcare/welcome>), with a final check for additional data sources using the Institute for Education Sciences website (<http://nces.ed.gov/pubsearch/>). We additionally limited our search to high-quality randomized controlled trials, considering only studies in the WWC that “meet evidence standards without reservations,” and we applied the same type of strict criteria for the design of randomized controlled trials when evaluating other potential study options on the additional websites.

The search for a recent randomized controlled trial study of early childhood education resulted in a list of 16 potential studies. However, several issues narrowed this list. First, some randomized controlled trials evaluated childcare subsidies or funding for preschool programs. Although these interventions included academic outcome measures, they are less likely to be interventions existing practitioners could easily implement in a different setting. Second, a number of studies focused on nonacademic outcomes such as social-emotional development. We chose to turn our attention to studies that included academically oriented intervention programs. Among these studies, a number were focused on subgroups of students, such as students with specific disabilities. Population data on students with specific disabilities can be more difficult to locate (although the Pre-Elementary Education Longitudinal Study does provide a nationally representative sample of children with disabilities).

In the end, we narrowed our list to approximately six studies that fit our criteria of high-quality randomized controlled trials of early childhood academic interventions focused on literacy or mathematics. Of these six studies, only one, Project Upgrade, had open access, publicly available data.

Project Upgrade, a 2-year randomized controlled trial, tested the effectiveness of three different language and literacy interventions in

childcare centers in Miami-Dade County, FL. We first examined this data set as a potential case study. The success of the reweighting approach relies on being able to adjust for a relatively large set of covariates that may differ between the trial sample and population and that moderate treatment effects. However, it proved difficult to identify many such covariates between Project Upgrade and a target population data set. We compared Project Upgrade to the HSIS (U.S. Department of Health and Human Services, 2002–2006), since both target low-income student populations. However, we found that although on the surface HSIS and Project Upgrade measured similar covariates for children, the studies utilize different specific measures of these covariates at both the classroom level and the individual level. For example, the HSIS classroom observation tool includes the Early Childhood Environment Rating Scale–Revised (ECERS-R), a rating scale that measures the quality of the classroom environment. In comparison, Project Upgrade utilizes a measure called Observation Measures of Language and Literacy Instruction in Early Childhood Education Classrooms (OMLIT). This battery of measures includes a series of questions related to the classroom environment, but the variables are measured differently than those included in the ECERS-R scale. In fact, Project Upgrade documentation includes a footnote indicating that they considered using the ECERS-R and rejected it in favor of the OMLIT, which they perceive to be a better measure of classroom environment and time use related to early childhood literacy. As a result, the two studies have topically similar assessments but a very limited number of common covariates. The two data sets did contain one identical scale, the Arnett Caregiver Rating Scale. However, even though the measure was the same, the Project Upgrade open access data file had the scale score variables already generated, rather than including the individual items. These variables were standardized to have a mean of 0 and a standard deviation of 1, but we believe this standardization used the mean for the Project Upgrade sample, rather than a national mean. As a result, when we generated a standardized score within the HSIS, the two standardized variables were not comparable, regardless of the fact that they had the same standardized mean and standard deviation.

Given the lack of sufficient covariate overlap between the Project Upgrade data and the HSIS, we turned back to the results of our initial search for early childhood randomized controlled trials, examining the five remaining randomized controlled trial studies of early childhood literacy or math. None of these studies had data that were readily available to researchers, so we contacted the principal investigator (PI) for the evaluation of the REDI intervention in Head Start centers to request data access. We were given permission to use requested parts of the data for a generalizability assessment and coordinated with the PI to receive the necessary data files.

Appendix B

Covariate Definitions and Comparability

This appendix details the creation of the eight common covariates between the REDI and HSIS samples and how the variables in each data set were made comparable.

- **Male:** This variable was generated from a single variable for child's gender in each of the two data sets.
- **Race dummy variables:** We created common race dummy variables by using multiple race variables in REDI to match the mutually exclusive categorical child's race variable that was already generated in the HSIS. The HSIS variable was broken into three categories: White/other, Black, and Hispanic. The REDI data contained survey responses that were not mutually exclusive categories; instead, the parents "marked all that apply" when indicating their child's race. We created a common set of race dummy variables by generating REDI variables to match the HSIS variables. All the White and Black respondents who indicated that they were Hispanic and Latino in the REDI data were coded as Hispanic. All non-Hispanic Black respondents were categorized as Black, and finally, all non-Hispanic White, Asian, and Other respondents were grouped in a White/Other category.
- **Spanish speaking:** Both surveys ask the parent whether the student spoke Spanish at home; however, the REDI survey asks the more general question of whether Spanish is a language used at home, while the HSIS survey asks the more specific question of which language is the primary language spoken at home. Thus, there may

be students who speak Spanish in the HSIS data set, but if they do not use it as their primary language at home, they may not be captured as Spanish speaking in this variable. Although this variable was generated as a potential covariate, the variable was ultimately not used in the propensity score model because of its collinearity with the Hispanic race/ethnicity dummy variable.

- Household size: This common variable was created from two variables in each data set, a variable counting children living in the household and a variable counting adults living in the household. In both data sets, the variables included the focal child and the survey respondent.
- Mother's marital status: The HSIS survey includes a question about maternal marital status. The REDI data set asked about the respondent's marital status. In most instances, the REDI respondent was the child's mother, but this was not universally true. So, to generate an equivalent measure in the REDI data, we combined several variables. The adult respondent was asked about their marital status, so if the respondent to the REDI survey was female, related to the child as a parent, and indicated that she was married, we combined those three variables to indicate that the child's mother was married.
- Mother's education: The HSIS data had a specific measure of child's mother's highest level of education. In the REDI data set, we used the highest level of education indicated by the respondent if she was female and related to the child as a parent to generate mother's highest level of education.
- Applied standard score: This is the Applied Problems mathematics test from the Woodcock Johnson III that assesses quantitative reasoning, but the test relies on language comprehension. Both data sets included this measure as a single variable in both the baseline data and postintervention wave of data collection. We utilize the baseline measure for both data sets.
- Elision score: Both data sets include a raw score for the Elision measure as a pre- and posttest.

Authors' Note

The content is solely the responsibility of the authors and does not necessarily represent the official views of the National Science Foundation, the National Institutes of Health, or the U.S. Department of Health and Human Services. The

information in this article was originally presented as part of a daylong symposium on external validity sponsored by the U.S. Department of Health and Human Services.

Acknowledgments

The authors particularly thank Dr. Karen Bierman for her assistance in obtaining the REDI data, and Cyrus Ebnesajjad for research assistance and the guest editor, Dr. T'Pring Westbrook, for organizing this special issue.

Declaration of Conflicting Interests

The author(s) declared no potential conflicts of interest with respect to the research, authorship, and/or publication of this article.

Funding

The author(s) disclosed receipt of the following financial support for the research, authorship, and/or publication of this article: Supported in part by Award DRL-1335843 from the National Science Foundation (co-PI's Stuart and Olsen) and by Award R305D150003 from the Institute of Education Sciences, U.S. Department of Education (co-PIs Stuart and Olsen). The original REDI study was supported as part of the Interagency School Readiness Consortium by National Institute of Child Health and Human Development grants HD046064 and HD43763. The HSIS was funded by the U.S. Department of Health and Human Services Administration for Children and Families, Office of Planning, Research and Evaluation, contract awarded to Westat, Inc. Contract # 282-00-0022.

Notes

1. Although the Head Start Impact Study (HSIS) sample was separated into treatment and control, we use the full sample of 4-year-olds, which is a nationally representative group of Head Start-eligible 4-year-old children. This is appropriate since no outcome data from the HSIS is used.
2. Note that the HSIS used an earlier but comparable version of this measure, which could be used in future work to do further diagnostics regarding the generalizability of the Research-Based, Developmentally Informed results to the HSIS.
3. For example, if the word is toothbrush, and you take away brush, what word does that make?
4. A measure of the household language being Spanish was also available in both sources but not able to be used because of collinearity with the indicator of Hispanic ethnicity.

References

- Bareinboim, E., & Pearl, J. (2013). A general algorithm for deciding transportability of experimental results. *Journal of Causal Inference*, 1, 107–134.
- Bauer, D. J., & Hussong, A. M. (2009). Psychometric approaches for developing commensurate measures across independent studies: Traditional and new models. *Psychological Methods*, 14, 101–125.
- Bell, S. H., Olsen, R. B., Orr, L. L., & Stuart, E. A. (2016). Estimates of external validity bias when impact evaluations select sites non-randomly. *Education Evaluation and Policy Analysis*, 38, 318–335.
- Bhattacharyya, A. (1943). On a measure of divergence between two statistical populations defined by their probability distributions. *Bulletin of the Calcutta Mathematical Society*, 35, 99–109.
- Bierman, K. L., Domitrovich, C. E., Nix, R. L., Gest, S. D., Welsh, J. A., Greenberg, M. T., . . . Gill, S. (2008). Promoting academic and social-emotional school readiness: The head start REDI program. *Child Development*, 79, 1802–1817.
- Braslow, J. T., Duan, N., Starks, S. L., Polo, A., Bromley, E., & Wells, K. B. (2005). Generalizability of studies on mental health treatment and outcomes, 1981–1996. *Psychiatric Services*, 56, 1261–1268.
- Bryant, P. E., MacLean, M., Bradley, L. L., & Crossland, J. (1990). Rhyme and alliteration, phoneme detection, and learning to read. *Developmental Psychology*, 26, 429.
- Catts, H. W., Fey, M. E., Zhang, X., & Tomblin, J. B. (1999). Language basis of reading and reading disabilities: Evidence from a longitudinal investigation. *Scientific Studies of Reading*, 3, 331–361.
- Cella, D., Yount, S., Rothrock, N., Gershon, R., Cook, K., Reeve, B., . . . On behalf of the PROMIS Cooperative Group. (2007). The Patient Reported Outcomes Measurement Information System (PROMIS): Progress of an NIH Roadmap Cooperative Group during its first two years. *Medical Care*, 45, S3–S11.
- Cole, S. R., & Stuart, E. A. (2010). Generalizing evidence from randomized clinical trials to target populations: The ACTG-320 trial. *American Journal of Epidemiology*, 172, 107–115.
- Cook, T. D. (2014). Generalizing causal knowledge in the policy sciences: External validity as a task of both multi-attribute representation and multi-attribute extrapolation. *Journal of Policy Analysis and Management*, 33, 527–536.
- Greenburg, D., & Schroder, M. (2004). *The digest of social experiments* (3rd ed.). Washington, DC: The Urban Institute Press.
- Hamilton, C. M., Strader, L. C., Pratt, J., Maiese, D., Hendershot, T., Kwok, R., . . . Haines, J. (2011). The PhenX toolkit: Get the most from your measures. *American Journal of Epidemiology*, 174, 253–260.

- Hirano, K., Imbens, G., & Ridder, G. (2003). Efficient estimation of average treatment effects using the estimated propensity score. *Econometrica*, 71, 1161–1189.
- Humphreys, K., Weingardt, K. R., & Harris, A. H. S. (2007). Influence of subject eligibility criteria on compliance with national institutes of health guidelines for inclusion of women, minorities, and children in treatment research. *Alcoholism: Clinical and Experimental Research*, 31, 988–995.
- Imai, K., King, G., & Stuart, E. A. (2008). Misunderstandings between experimentalists and observationalists about causal inference. *Journal of the Royal Statistical Society Series A*, 171, 481–502.
- Kent, D. M., Rothwell, P. M., Ioannidis, J. P. A., Altman, D. G., & Hayward, R. A. (2010). Assessing and reporting heterogeneity in treatment effects in clinical trials: A proposal. *Trials*, 11, 85.
- Kern, H. L., Stuart, E. A., Hill, J., & Green, D. P. (2016). Assessing methods for generalizing experimental impact estimate to target samples. *Journal of Research on Educational Effectiveness*, 9, 103–127.
- Lee, B. K., Lessler, J., & Stuart, E. A. (2010). Improving propensity score weighting using machine learning. *Statistics in Medicine*, 29, 337–346.
- McCaffrey, D. F., Ridgeway, G., & Morral, A. R. (2004). Propensity score estimation with boosted regression for evaluating causal effects in observational studies. *Psychological Methods*, 9, 403–425.
- Myers, J. A., Rassen, J. A., Gagne, J. J., Huybrechts, K. F., Schneeweiss, S., Rothman, K. J., . . . Glynn, R. J. (2011). Effects of adjusting for instrumental variables on bias and precision of effect estimates. *American Journal of Epidemiology*, 174, 1213–1222.
- Olsen, R., Bell, S., Orr, L., & Stuart, E. A. (2013). External validity in policy evaluations that choose sites purposively. *Journal of Policy Analysis and Management*, 32, 107–121.
- O’Muircheartaigh, C., & Hedges, L. V. (2014). Generalizing from unrepresentative experiments: A stratified propensity score approach. *Journal of the Royal Statistical Society: Series C*, 63, 195–210.
- Orr, L. L. (2015). 2014 Rossi Award Lecture: Beyond internal validity. *Evaluation Review*, 39, 167–178.
- Rothwell, P. M. (2005). External validity of randomised controlled trials: “To whom do the results of this trial apply?” *Lancet*, 365, 82–93.
- Rubin, D. B. (1974). Estimating causal effects of treatments in randomized and nonrandomized studies. *Journal of Educational Psychology*, 66, 688–701.
- Rubin, D. B., & Thomas, N. (1996). Matching using estimated propensity scores: Relating theory to practice. *Biometrics*, 52, 249–264.

- Schafer, J. L., & Kang, J. (2008). Average causal effects from nonrandomized studies: A practical guide and simulated example. *Psychological Methods, 13*, 279–313.
- Schochet, P. Z., Puma, M., & Deke, J. (2014). *Understanding variation in treatment effects in education impact evaluations: An overview of quantitative methods* (NCEE 2014-4017). Washington, DC: U.S. Department of Education, Institute of Education Sciences, National Center for Education Evaluation and Regional Assistance, Analytic Technical Assistance and Development. Retrieved from <http://ies.ed.gov/ncee/edlabs>
- Shadish, W. R., Cook, T. D., & Campbell, D. T. (2002). *Experimental and quasi-experimental designs for generalized causal inference*. Boston, MA: Houghton Mifflin Company.
- Stirman, S. W., Derubeis, R. J., Crits-Christoph, P., & Rothman, A. (2005). Can the randomized controlled trial literature generalize to nonrandomized patients? *Journal of Consulting and Clinical Psychology, 73*, 127–135. PMID: 15709839.
- Stuart, E. A. (2010). Matching methods for causal inference: A review and a look forward. *Statistical Science: A Review Journal of the Institute of Mathematical Statistics, 25*, 1–21.
- Stuart, E. A., Bradshaw, C. P., & Leaf, P. J. (2015). Assessing the generalizability of randomized trial results to target populations. *Prevention Science, 16*, 475–485. Retrieved from <http://www.ncbi.nlm.nih.gov/pmc/articles/PMC4359056/>
- Stuart, E. A., Cole, S., Bradshaw, C. P., & Leaf, P. J. (2011). The use of propensity scores to assess the generalizability of results from randomized trials. *Journal of the Royal Statistical Society Series A, 174*, 369–386.
- Stuart, E. A., Lee, B. K., & Leacy, F. P. (2013). Prognostic score-based balance measures can be a useful diagnostic for propensity score methods in comparative effectiveness research. *Journal of Clinical Epidemiology, 66*, S84–S90.
- Supplee, L. H., Kelly, B. C., MacKinnon, D. M., & Yoches Barofsky, M. (2013). Introduction to the special issue: Subgroup analysis in prevention and intervention research. *Prevention Science, 14*, 107–110.
- Susukida, R., Crum, R., Stuart, E. A., & Mojtabai, R. (2016). Assessing sample representativeness in randomized control trials: Application to the National Institute of Drug Abuse Clinical Trials Network. *Addiction, 111*, 1226–1234.
- Tipton, E. (2014). How generalizable is your experiment? An index for comparing experimental samples and populations. *Journal of Educational and Behavioral Statistics, 39*, 478–501.
- Tipton, E., Fellers, L., Caverly, S., Vaden-Kiernan, M., Borman, G., Sullivan, K., & Ruiz de Castilla, V. (In Press). Site selection in experiments: An assessment of site recruitment and generalizability in two scale-up studies. *Journal of Research on Educational Effectiveness*.

- Tipton, E., Hallberg, K., Hedges, L. V., & Chan, W. (In Press). Implications of small samples for generalization: Adjustments and rules of thumb. *Evaluation Review*.
- U.S. Department of Health and Human Services. Administration for Children and Families. (2002–2006). Office of Planning, Research and Evaluation. Head Start Impact Study (HSIS), 2002–2006 [United States]. ICPSR29462-v5. Ann Arbor, MI: Inter-university Consortium for Political and Social Research [distributor], 2014-03-21. doi:10.3886/ICPSR29462.v5
- U.S. Department of Health and Human Services, Administration for Children and Families. (2010, January). *Head Start Impact Study*. Final Report. Washington, DC: Author.
- Weintraub, S., Bauer, P. J., Zelazo, P. D., Wallner-Allen, K., Dikmen, S. S., Heaton, R. K., . . . Gershon, R. C. (2013). I. NIH toolbox cognition battery (CB): Introduction and pediatric data. *Monographs of the Society for Research in Child Development*, 78, 1–15. doi:10.1111/mono.12031
- Weiss, M. J., Bloom, H. S., & Brock, T. (2013). *A conceptual framework for studying the sources of variation in program effects*. MDRC Working Papers on Research Methodology. June 2013. Retrieved from http://www.mdrc.org/sites/default/files/a-conceptual_framework_for_studying_the_sources.pdf
- Wendling, B. J., Schrank, F. A., & Schmitt, A. J. (2007). *Educational interventions related to the Woodcock-Johnson III tests of achievement (Assessment service bulletin Number 8)*. Rolling Meadows, IL: Riverside Publishing.
- Westen, D. I., Stirman, S. W., & DeRubeis, R. J. (2006). Are research patients and clinical trials representative of clinical practice? In J. C. Norcross, L. E. Beutler, & R. F. Levant (Eds.), *Evidence-based practices in mental health: Debate and dialogue on the fundamental questions* (pp. 161–189). Washington, DC: American Psychological Association.

Author Biographies

Elizabeth A. Stuart is an associate dean for education and professor at Johns Hopkins Bloomberg School of Public Health, with appointments in Mental Health, Biostatistics, and Health Policy and Management. Trained as a statistician, her primary research areas are in methods for estimating causal effects, with applications in education, public health, and public policy.

Anna Rhodes is a doctoral student in the Sociology Department at Johns Hopkins University and a National Academy of Education/Spencer Dissertation Fellow. Her research examines how exposure to different educational and residential contexts affects students' outcomes.